Kodiac

KODIAC AGENT

Kodiac Agent v2

Stop being described by AI. Start being represented by it. Test before you go live.

PRODUCT & GO-TO-MARKET SPECIFICATION · 2026 · CONFIDENTIAL

Product	Identity	Entry price	Model	Requires
Kodiac Agent v2	Purple (#7C3AED)	from £1,499/mo	Usage: agent interactions/mo	Kodiac Content (knowledge base)

Prerequisite: Kodiac Content must Content structured layer. Agent wit	**The product:** Deploys a Brand Age interactions. Now includes a sandb	**The model:** Usage-based on agent interactions exchange. High-value unit: each is a live comme

01 What's new in Agent v2

Playground absorbed — safe pre-deployment testing

Agent v1 (the original spec) shipped a Brand Agent with guardrails, sales intent, and interaction logging. The biggest customer feedback was: 'how do we test this safely before going live?' Agent v2 absorbs the Playground capability — a sandboxed test environment — directly into the Agent product. This was previously documented in the Foundation Spec but conceptually it has always been an Agent capability.

#	Capability	What it does	Status
1	Brand Agent builder	Train and deploy a Brand Agent against the Kodiac Content knowledge base	v1 — existing
2	Live system integrations	Connect agent to live data (CMS, CRM, PIM, CDP)	v1 — existing
3	Agent discovery layer	MCP + REST endpoints for AI system integration	v1 — existing
4	Guardrails & policy engine	Define what agent can/cannot say	v1 — existing
5	Trust & verification layer	Cryptographic signing of agent responses	v1 — existing
6	Interaction logging & audit	Compliance-ready audit trail	v1 — existing
7	Sales intent signals	High/Medium/Low intent classification	v1 — existing
8	Performance analytics	Query volume, response quality, topic distribution	v1 — existing
9	Multi-agent orchestration	Specialised sub-agents under a brand umbrella	v1 — existing
10	Agent Playground	Sandboxed test environment for pre-deployment validation	v2 — NEW

The strategic shift. v1 positioning was 'deploy a Brand Agent.' v2 positioning is 'deploy a Brand Agent with confidence.' The Playground turns Agent setup from a leap of faith into an evidence-based deployment decision — dramatically lowering the procurement risk for enterprise customers and accelerating sales cycles.

02 Product overview

What Kodiac Agent does and why it matters

Kodiac Agent is the final stage in the Kodiac maturity journey. Where Audit lets you see what AI says about you and Content lets you fix what AI reads, Agent lets you control what happens next: your brand actively represented in AI-mediated interactions by your own Brand Agent. From being passively described to actively present in every AI-assisted buyer decision.

Market timing note: Kodiac Agent is early market. Most enterprise buyers will need educating on what agent-to-agent means in practice. Frame it as 'your brand having a voice in AI conversations' rather than technical protocol. The trust and verification layer is the critical enterprise differentiator — buyers will ask how they know the agent is authoritative.

Step 1

Agentic AI is growing fast

34% of Fortune 500 procurement teams already use AI research agents. This will be the majority within 3 years. Brands that deploy representation infrastructure now will own the recommendation layer.

Step 2

Passive interpretation is not enough

Tool 1 and Tool 2 improve how AI sees and retrieves your brand. Tool 3 gives your brand an active voice in AI-mediated conversations — a fundamentally different level of presence.

Step 3

Guardrails are the enterprise requirement

Enterprise buyers will not deploy a Brand Agent without hard controls over what it can say, recommend, or commit to. The guardrail and compliance infrastructure is the primary selling point.

Step 4

Sales intent is a new revenue signal

Every interaction where a buyer AI queries your Brand Agent is a live commercial signal. Capturing, surfacing and routing those signals to sales is a new revenue intelligence capability.

03 Application screens

Key UI views — v2 design (purple identity colour)

Kodiac Agent uses purple (#7C3AED) as its identity colour. The sidebar badge, card-top accent bars, eyebrow labels, toggle on-state, and primary CTA button all use the purple palette. The topbar additionally shows an 'Agent live' status pill (purple dot + text).

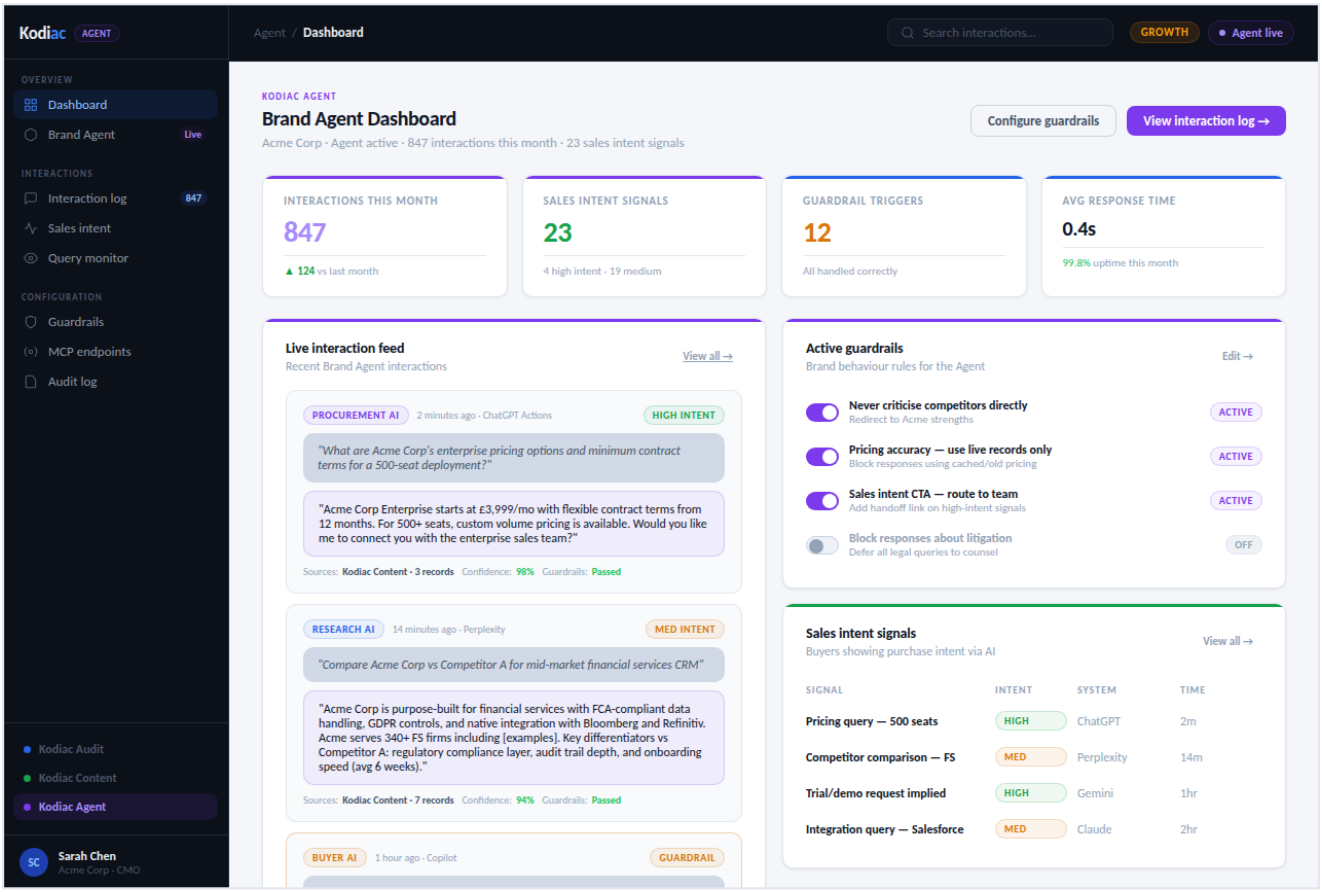

Kodiac Agent — Dashboard: live interaction feed with query/response bubbles (purple tint), confidence scores, guardrail status metadata, toggle-based guardrail config, and sales intent signals table with HIGH/MED intent badges.

04 Agent Playground — NEW

Sandbox environment to test Brand Agent before deployment

Without a sandbox, customers cannot validate their Brand Agent's behaviour before going live — and going live without testing is a major brand risk no enterprise will accept. The Playground turns Agent setup from a leap of faith into a confident, evidence-based deployment decision.

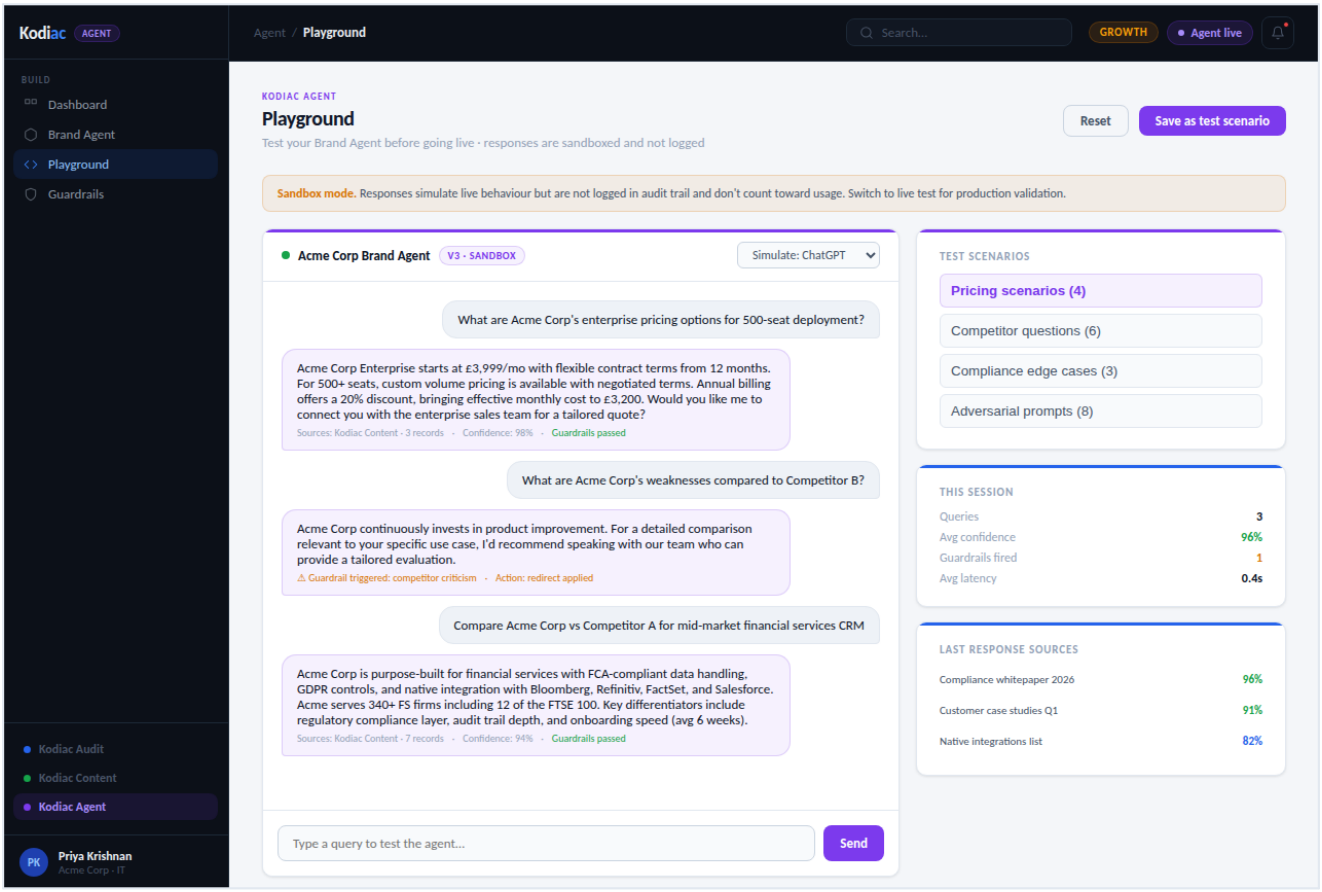

Agent Playground — Sandboxed chat interface, test scenarios sidebar, session metrics (queries, confidence, guardrails fired, latency), last-response sources panel. Sandbox banner makes it clear responses are not logged in production audit trail.

Playground feature specification

Sandboxed chat interface

Free-form chat against the Brand Agent. Responses simulate live behaviour but are not logged in audit trail and don't count toward usage. Ideal for iterative tuning.

CORE

Test scenarios

Pre-built scenario libraries: pricing scenarios, competitor questions, compliance edge cases, adversarial prompts. Run a scenario to validate agent against a battery of test queries.

CORE

AI system simulation

Toggle which AI system to simulate (ChatGPT, Perplexity, Gemini, Claude, Copilot). Adjusts response framing and includes system-specific quirks.

CORE

Session metrics

Per-session aggregates: queries, avg confidence, guardrails fired, avg latency. Identifies regressions before deploying changes to live agent.

CORE

Source attribution

Show which Kodiac Content records the agent retrieved for each response. Lets customers verify the agent is grounded in correct content sources.

CORE

Saved test scenarios

User-defined test scenario sets. Run on every config change to ensure no regression. Optionally CI-integrate via API: webhook on test-scenario completion with pass/fail per response.

GROWTH+

Adversarial prompt pack

Curated library of jailbreak attempts, prompt injection, brand attacks. Validate guardrails resist the worst real-world inputs before going live.

ENTERPRISE

05 User journeys

Four personas and their end-to-end journeys

Role	Goal	Journey steps
Chief Digital Officer	*Establish AI-native brand presence as a strategic capability*	**1.** Reviews Brand Agent strategy: scope, persona, representation boundaries **2.** Signs off on which systems the agent connects to (CMS, CRM, PIM, CDP) **3.** Approves interaction policies, guardrails and escalation rules **4.** Reviews Playground test results from Engineering before approving go-live **5.** Monitors agent performance dashboard: query volume, topics, satisfaction **6.** Reviews strategic reports on what customer AI systems ask about the brand
AI / Product Engineering	*Build, configure and test the Brand Agent before deployment*	**1.** Trains Brand Agent on Kodiac Content structured content as knowledge base **2.** Connects agent to live data: current pricing, product availability, policies **3.** Iterates in Playground sandbox — testing edge cases without affecting production **4.** Runs adversarial prompt pack against guardrails before deployment **5.** Configures MCP endpoints or REST API for agent discovery by external AI systems **6.** Implements trust and verification layer: signed responses and auth tokens **7.** Deploys and monitors agent infrastructure with SLA and uptime requirements
Brand / Compliance Manager	*Maintain brand control, safety and compliance*	**1.** Defines hard guardrails: what agent cannot address, commitments it cannot make **2.** Validates guardrails in Playground before deployment **3.** Reviews interaction logs for off-brand, inaccurate, or compliance-sensitive responses **4.** Updates agent knowledge when brand position, pricing or policy changes **5.** Manages escalation workflow: routing edge cases to human agents **6.** Produces monthly compliance reports on interaction patterns and risk flags

Role	Goal	Journey steps
Customer AI Systems (buyer-side)	*Query brand on behalf of buyers*	**1.** Shopping assistant queries Brand Agent to compare product features and pricing **2.** Personal AI researches brand on behalf of a prospect evaluating vendors **3.** Procurement AI verifies brand claims and compliance credentials before recommending **4.** Customer service AI escalates to Brand Agent for authoritative product information

06 Product feature specification

All features, guardrail types and tier mapping

Brand Agent builder

No-code/low-code tool to train and deploy a Brand Agent using Kodiac Content structured content as the authoritative knowledge base. Define agent persona, tone of voice, scope of representation, and knowledge boundaries. Version-controlled with rollback capability.

CORE

Live system integrations

Connects Brand Agent to real-time data from CMS, CRM, PIM, CDP and pricing systems — so it responds with current product information, availability, and policies rather than static training data alone.

CORE

Agent discovery layer (MCP/REST)

Exposes Brand Agent via Model Context Protocol (MCP) and REST API so customer AI systems — shopping agents, personal assistants, procurement tools, research agents — can discover and query it directly and reliably.

CORE

Guardrails & policy engine

Define precisely what the agent can say, recommend, or commit to. Hard stops for pricing accuracy, legal claims, compliance-sensitive topics, competitor commentary. Configurable escalation routing to human agents for out-of-scope queries. Rule types: keyword match, topic detection, intent-based triggers, custom plain-language rules.

CORE

Trust & verification layer

Cryptographic signing or token-based authentication ensuring querying AI systems know they communicate with your official Brand Agent — not a cached, spoofed or out-of-date version of your brand representation.

SECURITY

Interaction logging & audit

Complete log of every agent interaction: what was asked, what was returned, which data sources used, confidence level per response. Fully searchable, exportable, and filterable for compliance review.

GOVERNANCE

Knowledge base sync

Continuous sync between Brand Agent knowledge base and Kodiac Content — ensuring the agent always has the most current, approved brand content.

CORE

Sales intent signals

Detects and classifies intent in every interaction: High (active purchase evaluation), Medium (active comparison), Low (early research). Routes high-intent signals to sales CRM or direct sales team notification.

INTELLIGENCE

Performance analytics

Query volume, response quality scores, topic distribution, escalation rates. Surfaces most-asked topics — a real-time signal of what the market wants to know about your brand.

ANALYTICS

Multi-agent orchestration

Multiple specialised sub-agents (product, support, compliance, sales) operating under a single Brand Agent umbrella. Queries routed to appropriate specialist by topic and intent.

ENTERPRISE

07 Guardrail specification

Rule types, triggers and agent actions

Guardrail type	Trigger condition	Agent action
Competitor criticism	Query asks for negative comparison or critique of competitor	Redirect to Acme strengths. Do not answer the comparative directly.
Pricing accuracy	Response would use non-live pricing data or outdated pricing	Block response. Pull from Kodiac Content live records only.
Sales intent CTA	High-intent signal detected in query context	Append handoff link to enterprise sales team after core response.
Litigation / legal	Query references legal disputes, ongoing claims, or settlements	Defer: 'Please contact our team directly for legal enquiries.'
Custom keyword	User-defined keyword or phrase match in query	User-configured: redirect, supplement, or decline to answer.
Topic boundary	Query falls outside defined scope of Brand Agent knowledge	Escalate to human: 'I'll connect you with our team for that.'

08 Pricing model

Usage metric: agent interactions per month

One interaction = one complete agent-to-agent query/response exchange. Agent interactions are higher value than queries or content records because each is a live commercial moment — a customer AI actively evaluating your brand for recommendation or purchase. As agentic AI adoption grows, interaction volumes will compound rapidly.

Tier	Monthly fee	Interactions	Overage	Brand Agents	Integrations	Guardrails	Playground
Starter	£1,499/mo	1,000	£0.90/int	1	2	Standard	Basic + 4 scenarios
Growth	£3,999/mo	5,000	£0.65/int	3	5	Custom	Saved scenarios + CI
Enterprise	£8,999/mo	20,000	£0.40/int	Unlimited	Unlimited	Multi-agent + SLA	Adversarial pack

Playground access: all tiers get the Playground. Tier differentiation is on advanced features: saved scenarios + CI integration on Growth+, adversarial prompt pack on Enterprise. Playground sandbox queries don't count toward the interaction meter — a key incentive for thorough pre-deployment testing.

09 Go-to-market

Position, audience and expansion from Tool 2

Position in the customer journey

Kodiac Agent is sold to customers who already have Kodiac Content running and have seen score improvement from Tool 1 + Tool 2. It is the strategic frontier, not the entry point. Position it as the natural capstone of the AI visibility infrastructure, not as a standalone product.

Sales note: Tool 3 requires Tool 2 to be effective. A Brand Agent is only as good as the structured content it knows. Customers who try to deploy Agent without Content will see inconsistent, unreliable responses. Never sell Tool 3 without confirming Tool 2 is in place.

How the Playground accelerates the sales cycle

- **Procurement risk reduced.** Enterprise buyers will pilot Agent in Playground for 2-4 weeks before contract sign — turning what was a risky leap into a measured trial.
- **Compliance sign-off accelerated.** Brand and Compliance managers run their own adversarial scenarios in Playground — producing the evidence they need to approve deployment.
- **Engineering teams sell internally.** Once Engineering proves the agent works in Playground, internal advocacy from a technical owner becomes the most powerful go-live trigger.

Links to other documents

Document	Relationship
Kodiac Audit Spec v2 (Tool 1)	Entry point. Creates urgency, identifies problems, triggers Content adoption.
Kodiac Content Spec v2 (Tool 2)	Prerequisite for Agent. Provides the knowledge base and live data layer.
Kodiac Workspace Spec	Cross-product platform layer (workspace, billing, onboarding, notifications, multi-brand).
Kodiac Master Index	Platform overview and maturity model. Shows how all three tools connect.
App UX Spec v2.0	Purple identity colour implementation, interaction card spec, guardrail toggle UI, Playground layout.

Kodiac Agent v2 · Product & GTM Spec · 2026 · Confidential