Kodiac

KODIAC AUDIT

Kodiac Audit v2

See what AI sees, what AI reads, and what AI is told. The complete picture.

PRODUCT & GO-TO-MARKET SPECIFICATION · 2026 · CONFIDENTIAL

Product	Identity	Entry price	Model	Feeds into
Kodiac Audit v2	Blue (#2563EB)	from £499/mo	Usage: queries + pages + sources	Kodiac Content (fix engine)

The complete picture: Audit v2 me looks like to AI (input), and what AI	**The biggest gap closed:** 73% of A Wikipedia). The original spec meas	**Why this matters:** Customers can no longer fix see and prioritise interventions across the entire

01 What's new in Audit v2

Two new capabilities that complete the audit picture

Audit v1 (the original spec) measured one layer: what AI systems say about your brand when prompted (the “output layer”). Audit v2 adds two more measurement layers (Website Score, Source Intelligence) and absorbs two configuration capabilities (Prompt Library, Source of Truth) that were previously documented in the Foundation Spec. The result is a complete, self-contained Audit product.

#	Capability	What it measures / does	Status
1	AI Output Audit	What ChatGPT, Perplexity, Gemini, Claude, Copilot say about you	v1 — existing
2	Website Score	What AI sees when crawling your owned web pages — 10-dimension AI-readiness score	v2 — NEW
3	Source Intelligence	Third-party content AI reads about you (Reddit, G2, Wikipedia, news, forums) and how it shapes responses	v2 — NEW
4	Prompt Library	Manage the prompts Audit sends to AI systems. Token expansion, versioning, generated starter set.	v2 — absorbed from Foundation
5	Source of Truth	Brand fact registry powering inaccuracy verification. Auto-syncs from Content where connected.	v2 — absorbed from Foundation

The strategic shift. v1 positioning was “see what AI says.” v2 positioning is “see, diagnose at source, and intervene across the entire information ecosystem about your brand.” This is a much stronger and more defensible category position. Customers cannot fix poor AI scores by improving their website alone — they need to address every layer that shapes AI representation. Audit v2 is what makes that possible.

02 Prompt Library & Source of Truth

Two configuration capabilities, now formally part of Audit

These two capabilities were previously documented in the Foundation Spec but conceptually they have always been Audit capabilities. They configure what Audit asks AI systems and what Audit verifies responses against. Moving them into Audit v2 gives customers a complete, self-contained Tool 1 spec and clears the Foundation Spec to focus on cross-product platform capabilities.

Prompt Library

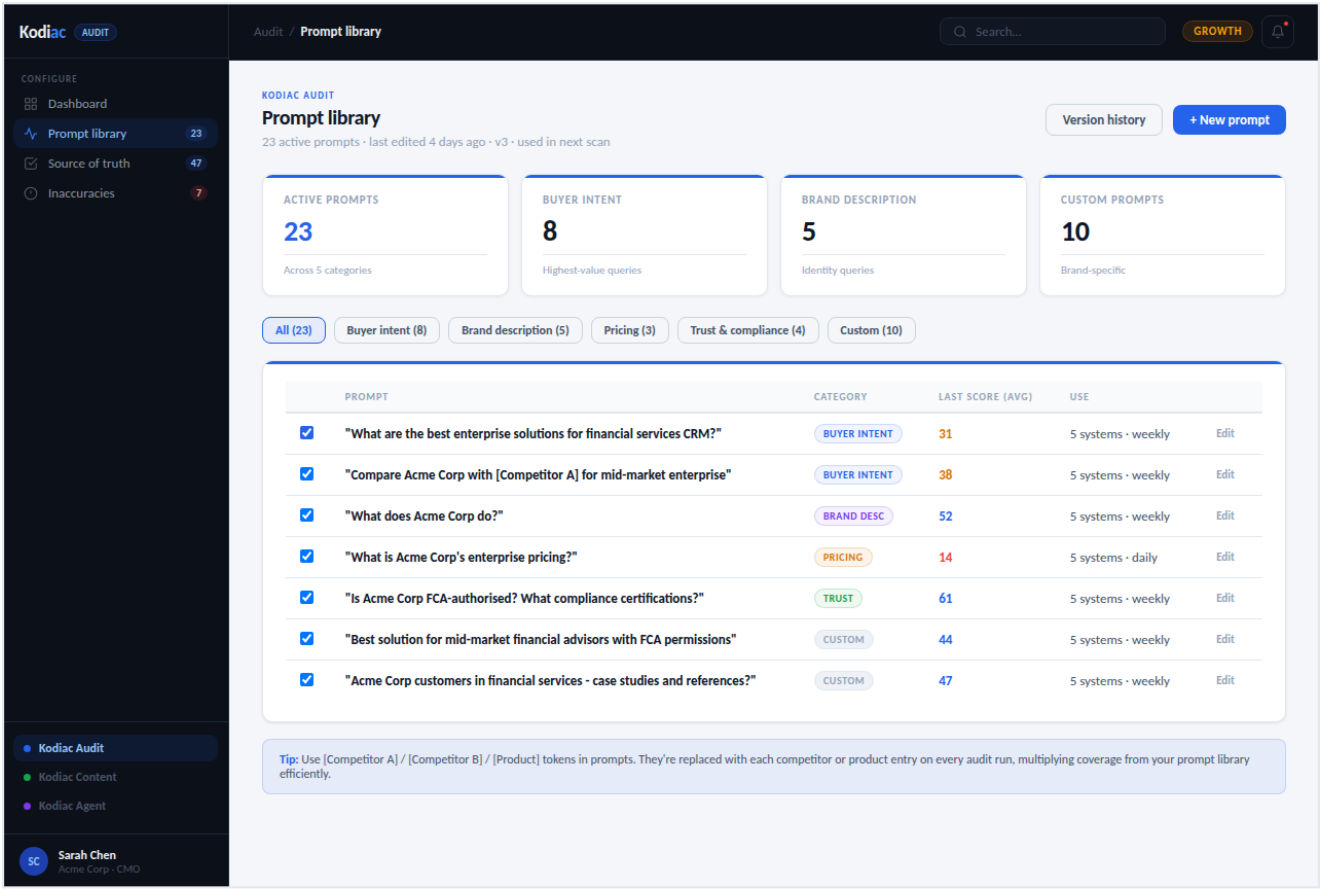

Prompt Library — 23 active prompts across 5 categories, last-score per prompt, filter pills, version history, [Competitor]/[Product] token system for efficient prompt expansion.

Generated starter library

On onboarding, generate 12 prompts based on industry, brand, and competitors. Categories: buyer intent, brand description, pricing, trust/compliance, custom.

CORE

Custom prompt authoring

Users add unlimited custom prompts (subject to plan tier query limits). Categorise and tag for filtering. Plain-language with token system: [Competitor], [Product], [Locale].

CORE

Token expansion

[Competitor] expands to each of your tracked competitors per audit run. 12 prompts × 5 competitors = 60 effective queries per run. Maximises prompt library value.

CORE

Versioning

Every prompt-library save creates a version. Diff between versions. Roll back if a change degraded results. Tag versions (e.g. 'v3 — added pricing prompts').

CORE

Per-prompt analytics

Last score, trend per prompt, per-AI-system breakdown. Surface which prompts produce the worst results — those are usually root cause indicators.

CORE

Industry prompt packs

Curated industry-specific prompt packs (Financial Services, Healthcare, Retail, B2B SaaS). One-click import. Maintained by Kodiac team.

GROWTH+

Source of Truth

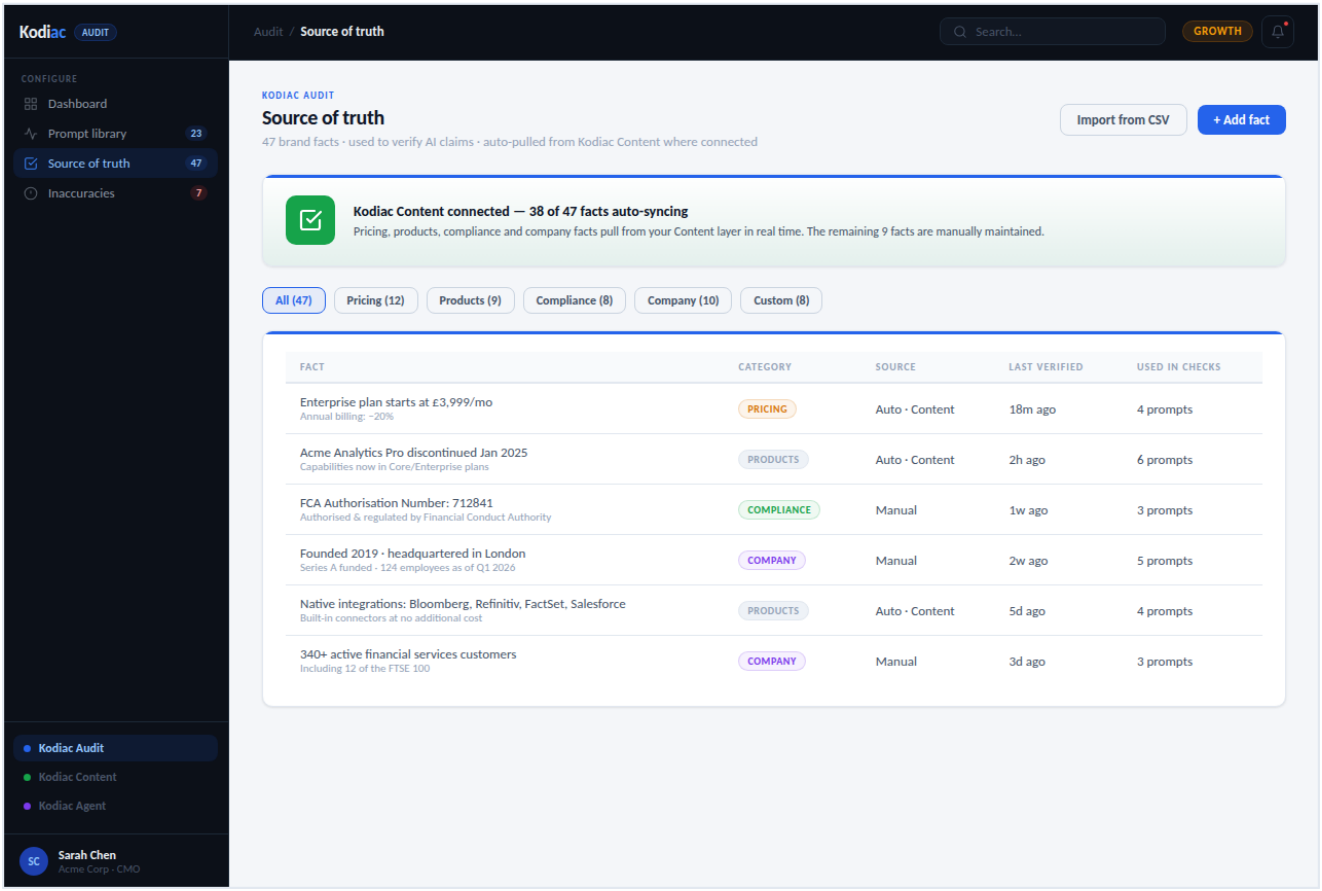

Source of Truth — 47 brand facts across 5 categories. Shows when facts auto-sync from Kodiac Content vs manually maintained. Per-fact: source, last verified, prompts using it.

Manual fact registry

Up to 50 facts on Starter, 200 on Growth, unlimited on Enterprise. Each fact: name, category, value, source (manual or auto), effective dates, used-in-prompts list.

CORE

Auto-sync from Content

When Kodiac Content is connected, pricing/products/compliance/company facts auto-pull from the live structured layer. No duplication, no drift between Tool 1 and Tool 2.

CORE

Inaccuracy verification

On every audit run, AI-claims are extracted and compared to source of truth. Mismatch triggers Critical/Moderate/Minor severity flag depending on category and deviation.

CORE

Fact freshness

Each fact has a 'last verified' timestamp. Stale facts (>90 days) flagged for review. Auto-pulled facts inherit freshness from Content sync time.

CORE

Bulk import

CSV upload for initial setup. Mapping wizard handles arbitrary CSV structures. Re-import to bulk update without overwriting manually entered facts.

CORE

04 Product overview

Three-layer audit — output, input, ecosystem

Kodiac Audit v2 is the entry point to the Kodiac platform. It runs three complementary scans to give customers a complete, actionable picture of how AI systems represent their brand.

The three-layer audit

Step 1

AI Output Audit

Sends brand-relevant prompts to 5 AI systems. Captures, scores and tracks responses. Detects inaccuracies. Benchmarks vs competitors.

Step 2

Website Score

Crawls your owned web pages. Scores AI-readiness across 10 dimensions including structured data, semantic clarity, factual density, and crawlability.

Step 3

Source Intelligence

Discovers third-party content AI systems read about your brand. Categorises, scores AI weight, and recommends prioritised interventions.

Why all three layers matter

- **Output without input is theatre.** Knowing AI describes you poorly without diagnosing why creates urgency without action.
- **Input without ecosystem is incomplete.** Improving your website fixes 27% of typical AI representation issues. The other 73% comes from third-party sources.
- **Ecosystem without output is unfocused.** Without knowing which AI responses are poor, ecosystem interventions can't be prioritised.
- **All three together = a complete diagnostic.** Customer sees the symptom (output), the source-of-truth gap (input), and the third-party narrative (ecosystem) — and can act on each independently.

04 Website Score — NEW

AI-readiness diagnostic for your owned web pages

Website Score crawls your domain and rates each page across 10 dimensions of AI-readiness. It produces a top-line score (0-100), a per-dimension breakdown, a per-page breakdown showing the lowest-scoring pages, and a fix list directly tied to the inaccuracies and gap analysis from the AI Output Audit.

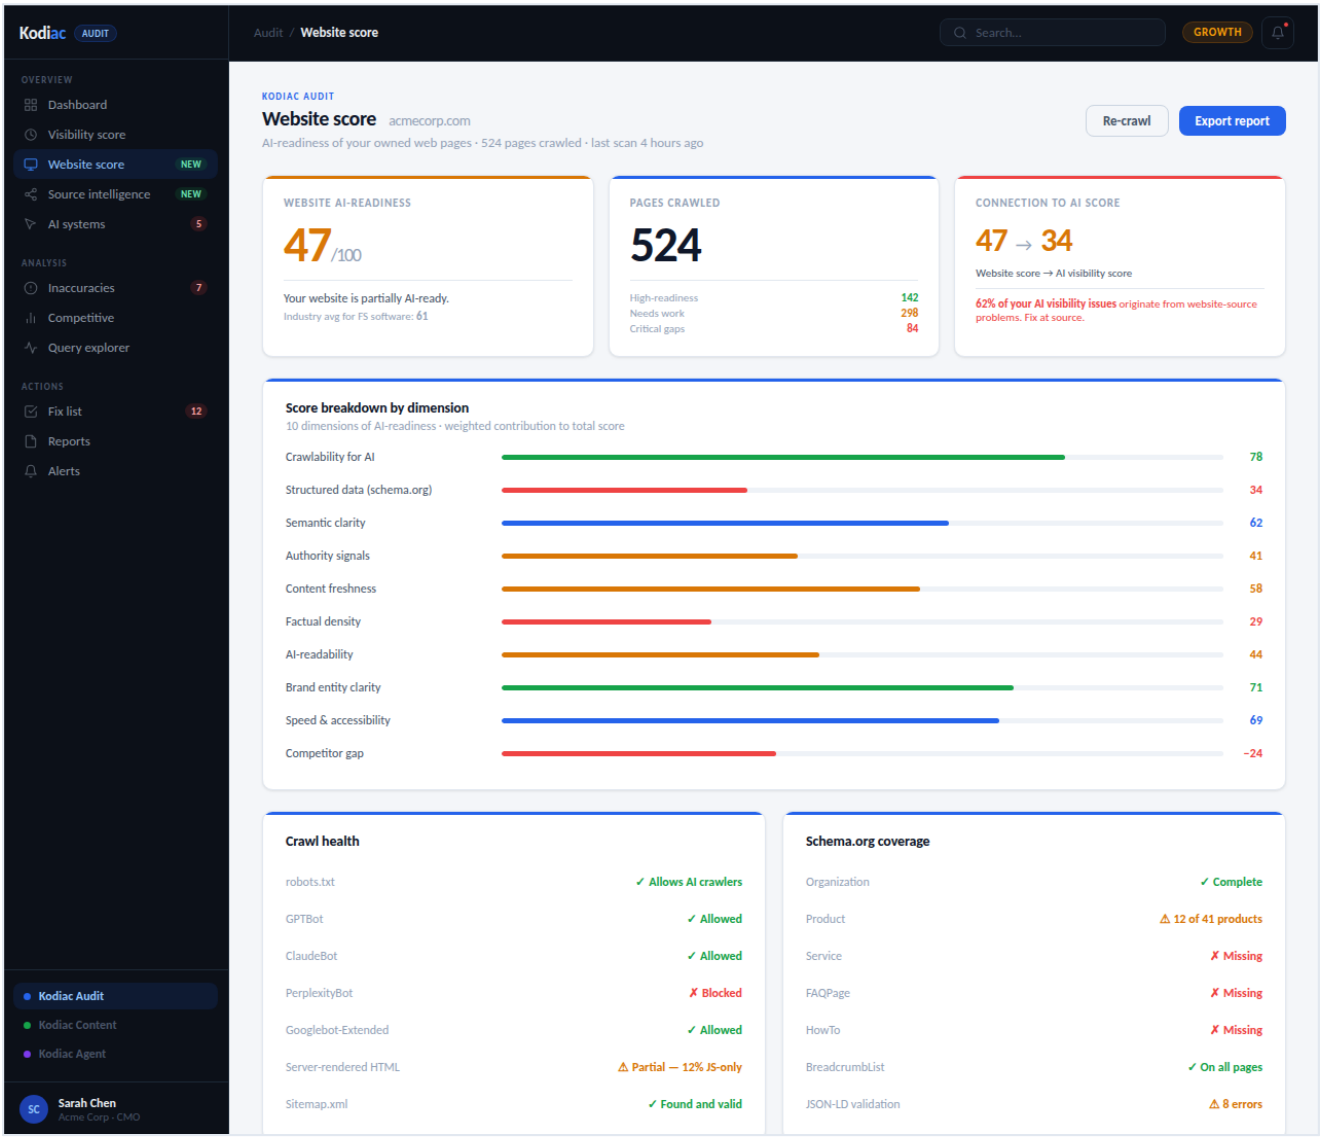

Website Score — Top-line score (47/100), 10-dimension breakdown, crawl health, schema.org coverage, lowest-scoring pages with per-page Fix CTAs.

The 10 dimensions of AI-readiness

#	Dimension	What it measures	Why it matters to AI
1	Crawlability for AI	robots.txt allows AI crawlers (GPTBot, ClaudeBot, PerplexityBot, Googlebot-Extended). Server-rendered HTML available. JS-only content minimised.	If AI can't crawl, AI can't represent.
2	Structured data	schema.org coverage (Organization, Product, Service, FAQPage, HowTo, BreadcrumbList). JSON-LD validation. Field completeness.	Structured data is the AI training signal of choice.
3	Semantic clarity	Clear h1/h2 hierarchy. Descriptive titles. Meta descriptions. Alt text. Semantic HTML5 elements (article, section, nav).	Helps AI understand what each page is about.
4	Authority signals	About page completeness. Author bylines. Citations. Inbound links from authoritative domains. Wikipedia/Wikidata presence.	Authority drives weighting in AI training corpora.
5	Content freshness	Last-updated timestamps. Content age distribution. Dead pages. Broken links. Update cadence.	Stale content is deprioritised in AI retrieval.
6	Factual density	Specific facts AI can extract: prices, dates, dimensions, qualifications, certifications, numbers. Vague marketing prose flagged.	AI prefers extractable facts over hand-waving.
7	AI-readability	Content chunked into AI-extractable pieces. FAQ patterns. Definition patterns. Comparison tables. Step-by-step formats.	Chunkable content retrieves cleanly into RAG pipelines.
8	Brand entity clarity	Consistent brand naming. Knowledge graph presence. Social profile links. Structured contact info. Sameas pointers.	AI builds a coherent entity model of your brand.
9	Speed & accessibility	Core Web Vitals. Accessibility score. Mobile rendering. Resource size budgets.	Performance metrics carry forward as AI crawl signals.
10	Competitor gap	Same metrics computed for top-3 named competitors. Side-by-side comparison.	Tells customers where they're falling behind specifically.

Website Score features

Full domain crawl

Discovers all pages on the customer's domain via sitemap.xml, links, and robots.txt. Configurable depth and exclude patterns. Respects crawl-delay. Crawls 10K pages on Starter, 50K on Growth, 250K on Enterprise per audit cycle.

CORE

Per-page scoring

Each page receives its own 0-100 score across all 10 dimensions. Low-scoring pages flagged for immediate attention. Bulk-update mode: select pages, push to Content fix list, fix once.

CORE

AI crawler detection

Identifies which AI crawlers (GPTBot, ClaudeBot, PerplexityBot, Googlebot-Extended, etc.) are allowed/blocked. Surfaces unintentional blocks. Recommends optimal robots.txt configuration.

CORE

Schema.org coverage analysis

Maps every structured data type present, missing, or invalid. JSON-LD syntax validation. Highlights high-impact missing types (FAQPage, HowTo for B2B SaaS; Product/Offer for ecommerce).

CORE

Connection-to-output mapping

Algorithmically links website-score gaps to specific AI Output Audit findings: “This pricing page issue causes Perplexity to misquote your prices in 4 of 12 buyer queries.”

INTELLIGENCE

Per-page fix CTAs

Each low-scoring page has a Fix button that creates a record in Content's fix list (if Content is connected) or a CSV export of fixes if not.

CORE

Continuous re-crawl

Weekly (Starter), daily (Growth), real-time on URL change (Enterprise). Drift detection: alert if a page's score drops by 10+ points.

MONITORING

Competitor crawl

Same crawl and scoring run on competitor domains for comparison. 1 competitor (Starter), 5 (Growth), 20 (Enterprise).

INTELLIGENCE

API & webhooks

REST API for crawl-on-demand, webhook on score change, bulk export. CI integration possible for marketing teams running tests in staging.

ENTERPRISE

05 Source Intelligence — NEW

Third-party content AI reads about your brand

Source Intelligence identifies the third-party content sources AI systems are using when generating responses about your brand. It categorises sources (Reddit, G2, Wikipedia, news, forums, YouTube transcripts, Hacker News, etc.), scores their AI weight, evaluates sentiment, and recommends the highest-leverage interventions — prioritised by predicted score uplift.

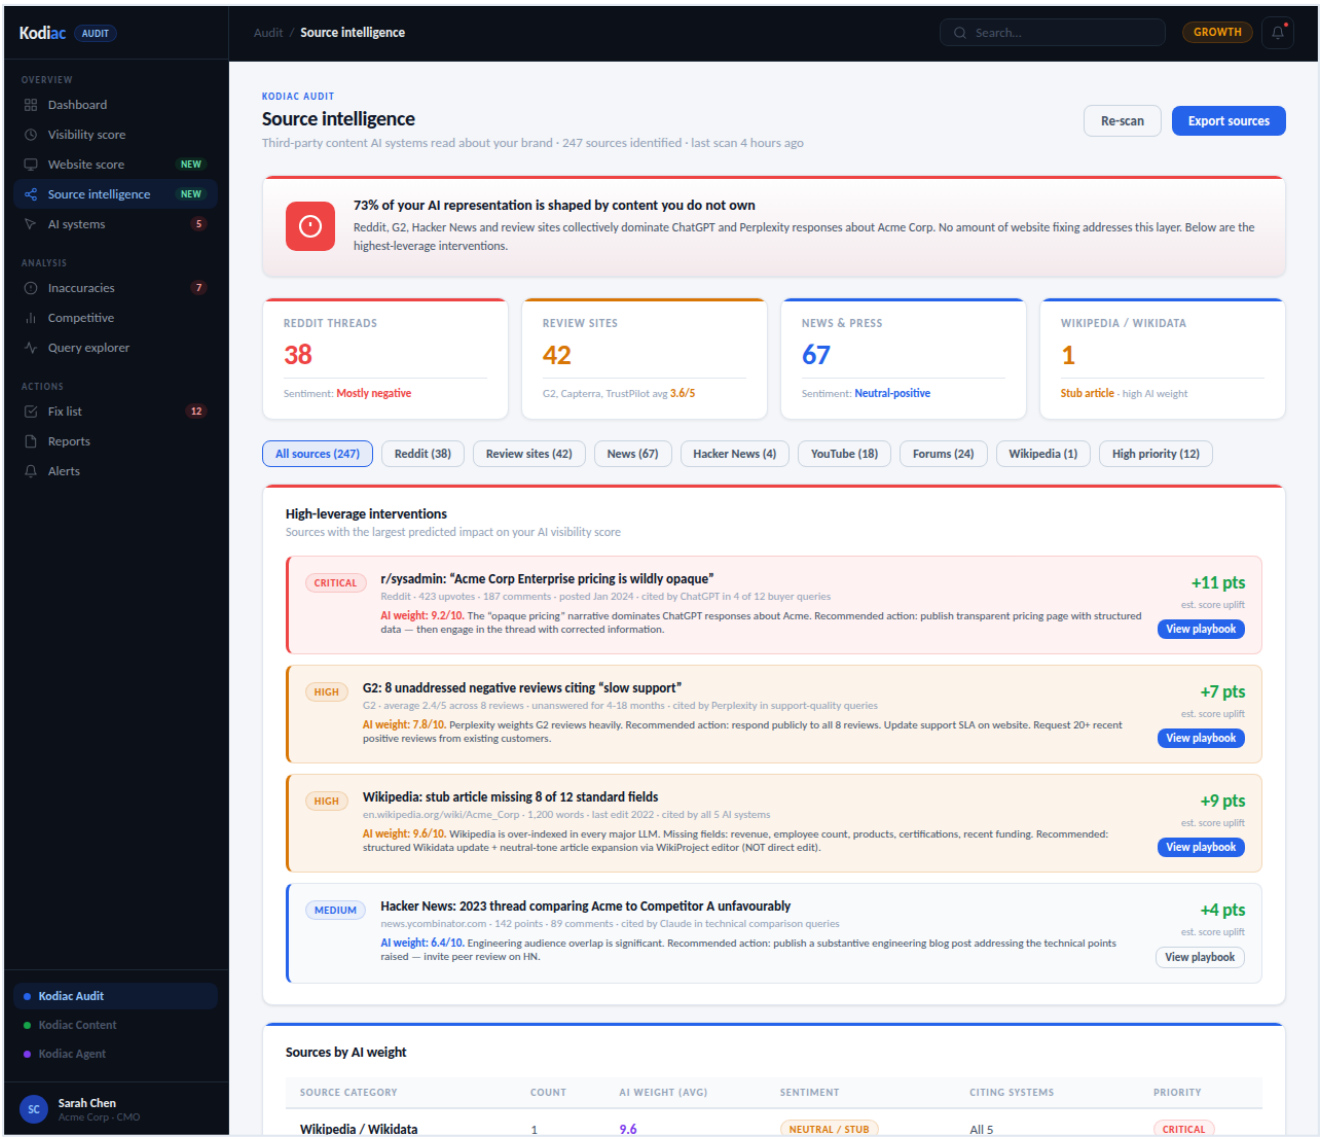

Source Intelligence — Strategic context callout, source breakdown by category (Reddit, reviews, news, Wikipedia), high-leverage interventions list with playbooks, sources-by-AI-weight table.

Source categories monitored

Category	Examples	Typical AI weight	Why it matters
Wikipedia / Wikidata	en.wikipedia.org/wiki/[brand]	9.0–10.0	Most over-indexed source in every major LLM. One stub article can dominate all AI responses.
Reddit	r/sysadmin, r/ecommerce, r/finance, sector-specific subs	8.0–9.5	Heavily over-indexed in ChatGPT and Perplexity. Long-tail negative threads can poison responses for years.
Review sites	G2, Capterra, TrustPilot, Glassdoor	7.0–8.5	Treated as semi-authoritative. Perplexity weights heavily. Negative trends compound.
Hacker News	news.ycombinator.com	6.0–7.5	Disproportionate weight for technical/B2B brands. One thread can shape engineering audience perception.
News & trade press	FT, WSJ, Bloomberg, sector publications	5.5–7.0	Established sources, weighted high. Often the source of factual claims in AI output.
YouTube transcripts	Review channels, conference talks, podcasts	5.0–6.5	Indexed and transcribed at scale. Influences sentiment and technical understanding.
Industry forums	Sector-specific communities (e.g. r/financialcareers, advicelocal forums)	4.5–6.0	Niche but trusted within vertical. Procurement teams research here.
Stack Overflow / GitHub	For technical brands	4.0–5.5	Dominate Claude's view of technical brands. Issue threads matter.
Quora / Medium / Substack	Long-tail opinion content	3.0–4.5	Volume matters — dozens of low-weight sources can compound.
Glassdoor	Employee reviews	3.5–5.0	Influences talent and culture queries. Some buyer-side weight.

Source Intelligence features

Source discovery via AI response analysis

Runs the prompt library through 5 AI systems and analyses cited sources directly (when AI shows them, e.g. Perplexity, Gemini) and infers them from response patterns (for ChatGPT, Claude when sources aren't shown). Combined view across all 5 systems.

CORE

AI weight scoring

Each source is scored 0-10 on its AI weight — i.e. how much that specific source is shaping AI's view of your brand. Computed from citation frequency, source authority in AI training data, and response-content overlap.

CORE

Sentiment analysis

Per-source sentiment classification (positive / neutral / mixed / negative). Tracked over time. Sentiment trend alerts (e.g. “Reddit sentiment dropped to negative this week”).

CORE

Predicted score uplift

Each high-priority source is scored on predicted AI visibility score uplift if addressed. Ranks interventions by expected impact, not by source size.

INTELLIGENCE

Intervention playbooks

Per source category, a tailored playbook of recommended actions: (a) what Kodiac can automate (Wikidata structured updates, schema submissions to G2), (b) what requires customer action (Reddit engagement, Wikipedia editor outreach via WikiProject), (c) what requires content marketing (HN-tier engineering blog post, YouTube creator outreach).

CORE

Continuous monitoring

Weekly (Starter), daily (Growth), real-time alert for new high-weight sources (Enterprise). New negative source triggers a critical alert. Sentiment shifts trigger warnings.

MONITORING

Wikipedia / Wikidata sub-module

Special handling for Wikipedia given outsized AI weight. Identifies missing fields, stub status, last-edit recency, edit-conflict risk. Recommends WikiProject engagement (NOT direct edits, which can backfire on neutral-tone violations).

INTELLIGENCE

Review-site sub-module

G2/Capterra/TrustPilot integration. Identifies unanswered negative reviews, response-rate, recency. Suggests review-request workflows for happy customers.

INTELLIGENCE

Source change alerts

New high-weight source appears → alert. Existing source sentiment shifts → alert. Negative thread gains traction → alert. Configurable per category.

MONITORING

Source export & integrations

CSV/JSON export of all sources. Webhook on new source. Push to PR/comms team via Slack/email/Teams. CRM integration for sales-relevant sources.

ENTERPRISE

06 User journeys

Updated personas — now using all three audit layers

Role	Goal	Journey steps
Chief Marketing Officer	*Understand and report AI brand exposure risk — across the whole information ecosystem*	**1.** Receives audit and sees alarming AI Output score (34/100) on Dashboard **2.** Drills into Website Score — 47/100 — sees 62% of AI issues stem from website-source problems **3.** Drills into Source Intelligence — sees 73% of AI representation comes from third-party sources **4.** Reviews highest-leverage interventions: Wikipedia stub (+9 pts), Reddit threads (+11 pts), G2 reviews (+7 pts) **5.** Exports board report showing the full diagnosis: output, input, and ecosystem
Head of Digital / SEO	*Diagnose root causes across all three layers and execute fix list*	**1.** Notices critical inaccuracies on Audit Dashboard **2.** Drills into Website Score: identifies 84 critical-gap pages, mostly missing structured data **3.** Cross-references with Source Intelligence: confirms low Wikipedia presence + negative Reddit narrative **4.** Builds a 30/60/90-day plan: (1) fix top-20 pages, (2) Wikidata update, (3) review platform engagement **5.** Tracks improvement weekly across all three layers in unified dashboard
Brand / Comms Manager	*Monitor brand narrative across the entire third-party ecosystem*	**1.** Sets alert rules: new negative Reddit thread, sentiment drop on G2, new news article **2.** Receives critical alert: r/sysadmin thread reaches AI weight 9+ **3.** Reviews recommended playbook: don't engage hostilely, publish corrected pricing page first, then comment **4.** Coordinates response with PR team via Slack integration **5.** Tracks sentiment recovery over following weeks
Agency / Consultant	*Deliver three-layer audit as managed service*	**1.** Runs three-layer audit on prospect brand pre-pitch — result: AI=28, Website=39, Sources=23 (Reddit dominated) **2.** Walks prospect through findings: “your AI score isn't a website problem; it's a Reddit problem” **3.** Pitches three-layer managed service: £X/month for ongoing monitoring + interventions **4.** Wins client. Configures three-layer monitoring on signing. **5.** Delivers QBR showing improvement across output, input, and ecosystem layers

07 Pricing model v2

Three usage meters: queries, pages, sources

Audit v2 introduces two additional usage meters alongside the existing brand queries meter. Pages crawled = Website Score volume. Sources monitored = Source Intelligence volume. Each scales independently with the value the customer derives. Tier limits remain at the same headline price.

Tier	Monthly fee	Brand queries	Pages crawled	Sources monitored	AI systems	Competitors	Cadence
Starter	£499/mo	500 / mo	10,000 / cycle	100	5	2	Weekly
Growth	£1,499/mo	2,500 / mo	50,000 / cycle	500	All major	10	Daily
Enterprise	£3,999/mo	10,000 / mo	250,000 / cycle	Unlimited	All + custom	Unlimited	Real-time

Overage rates

Meter	Starter overage	Growth overage	Enterprise overage
Brand queries	£0.60/query	£0.45/query	£0.30/query
Pages crawled	£2 / 1,000 pages	£1.50 / 1,000 pages	£1 / 1,000 pages
Sources monitored	£1 / source / month	£0.75 / source / month	£0.50 / source / month

Pricing principle preserved. No feature gates between tiers — all tiers access the full feature set (AI Output, Website Score, Source Intelligence). Tiers differ only in volume allowances, cadence, and support level. Annual billing: –20% on all tiers. Existing v1 customers: free upgrade to v2 with same headline price. v2 launches with 30-day grace period before overage on new meters takes effect.

08 GTM impact of v2

How Audit v2 changes positioning, sales motion, and expansion

Positioning shift

Dimension	v1 positioning	v2 positioning	Why v2 is stronger
Headline	See what AI says about your brand	See, diagnose at source, and intervene across the entire information ecosystem about your brand	Bigger and more defensible category position
Tagline	“AI is recommending your competitors. Find out why.”	“AI is shaped by content you don't own. Take control of all of it.”	Captures the third-party reality customers will increasingly recognise
Free wedge	Free 60-second AI audit	Free three-layer audit (AI + website + top 5 sources)	More compelling free experience drives higher conversion

Sales motion impact

- **Outbound emails get sharper.** “We ran your audit — 73% of how AI describes you comes from a Reddit thread you don't know about” outperforms generic AI-visibility outreach.
- **Pitch meetings have a stronger story.** Three-layer findings make it harder for the prospect to dismiss the diagnosis as a website-only problem they can fix internally.
- **Competitive sales motion is stronger.** Source Intelligence is genuinely first-to-market. No competitor (incumbent or startup) is doing this. Proprietary data and category-defining language.

Expansion mechanics

- **Audit → Content expansion strengthens.** Website Score gaps map directly to Content fix list. The cross-product workflow becomes seamless: low-scoring page → Fix CTA → Content record → publish to AI.
- **Source Intelligence creates new expansion conversations.** “Your Wikipedia presence is a stub — here's the playbook to fix that” is a different sale than “your content needs structured data.” Some of those conversations lead to managed-services upsell or partner-agency referral.
- **Volume meters create natural expansion.** Brands monitoring 100 sources at Starter who upgrade to Growth at 500 sources experience compounding overage. The scaling growth scenario in the financial model now has three independent levers, not one.

09 Cross-document references

How v2 changes related documents

Document	Update needed	Status
Kodiac Audit Spec (this document)	Replaces v1 spec entirely	✓ Done — supersedes Kodiac_Audit_Spec.pdf
Kodiac Master Index	Update Audit description from “what AI says” to “three-layer audit”. Update revenue model with three-meter pricing.	Pending
App UX Spec v2.0	Add Section 11 covering audit-website-score.html and audit-sources.html with screenshots	Pending
App Nav Flow doc	Add 2 new screens to catalogue. Add new use case: SEO/Digital lead diagnoses across all three layers.	Pending
Strategic Roadmap	Source Intelligence already in Year 2 strategic bets — actually, given v2 launches it sooner, update Year 1 to include Source Intelligence as launched in Q2 2026 alongside Content GA.	Pending
Internal Execution Plan	Update Q1 2026 product OKRs to include Audit v2 capabilities. Add new capability to Q2 2026 release schedule (Source Intelligence GA).	Pending
Marketing Site	Update audit/free-audit pages to reflect three-layer audit. Update pricing page with new meters.	Pending
LinkedIn Playbook	Add Source Intelligence to launch carousel concepts. New post template: “73% of your AI representation comes from sources you don't own.”	Pending
Investor Pitch	Update slide on product depth to include three-layer audit. Bigger TAM story.	Pending

Kodiac Audit v2 · Product & GTM Spec · 2026 · Confidential