Kodiac

KODIAC CONTENT

Kodiac Content v2

See, fix, and serve. The complete content layer for AI.

PRODUCT & GO-TO-MARKET SPECIFICATION · 2026 · CONFIDENTIAL

Product	Identity	Entry price	Model	Feeds into
Kodiac Content v2	Green (#16A34A)	from £799/mo	Usage: records + API calls	Kodiac Agent (knowledge base)

Prerequisite: Recommended after visibility scores. Can be adopted in	**The product:** Connects to your CM records inline, and serves the struc	**The model:** Usage-based on records managed expansion trigger into Kodiac Agent.

01 What's new in Content v2

Authoring and developer surface absorbed

Content v1 (the original spec) measured one job: ingest content from CMS connectors, normalise and govern it, publish to AI surfaces. v2 adds two more capabilities, completing the content layer.

#	Capability	What it does	Status
1	Universal content connectors	Ingest from 15+ source systems with delta sync	v1 — existing
2	Normalisation engine	Transform to unified AI-readable content model	v1 — existing
3	Structured content modelling	Define content models optimised for AI consumption	v1 — existing
4	Governance & approval layer	Status flags, approval workflows, audit trails	v1 — existing
5	Middleware delivery API	Content delivery API for downstream AI tools	v1 — existing
6	Schema.org & JSON-LD export	Structured data markup auto-generation	v1 — existing
7	Content quality scoring	Automated AI-readiness scoring per record	v1 — existing
8	Delta sync & versioning	Efficient incremental sync with version history	v1 — existing
9	Content authoring	Inline structured editor + write-back to source CMS	v2 — NEW
10	Developer surface	RAG endpoint, vector index, embedding config, live test query	v2 — NEW

The strategic shift. v1 positioning was 'AI-ready content delivery.' v2 positioning is 'the complete content layer for AI — see what AI reads, fix it inline, serve it as structured data and via vector retrieval.' This is a much stronger and more defensible category position. v2 also closes the Audit→Content fix loop tighter — a customer can fix a low-scoring page in Authoring and watch the AI visibility score recover.

02 Product overview

What Kodiac Content does and why it exists

Kodiac Content is the fix engine. Where Kodiac Audit identifies what AI is getting wrong about your brand, Content fixes it. It connects to every source system in your content estate, normalises and governs the content, and publishes it as structured, AI-readable data that AI systems can trust, retrieve accurately, and reuse. It is also the knowledge base that powers Kodiac Agent.

The problem it solves

Step 1

Fragmented content estate

Enterprise content is spread across 10+ systems: CMS, DAM, CRM, SharePoint, Dropbox. None is structured for AI retrieval.

Step 2

AI cannot trust your content

AI systems retrieve by semantic meaning. Your URL structure, CMS taxonomy and content governance have zero influence.

Step 3

No governance layer

Enterprise buyers need controls: which content AI can use, which is restricted, what has expired. None of their current tools provide this.

Step 4

Editing live in Kodiac matters

With v1, fixes had to happen in source CMS first. With v2, customers edit in Kodiac and write back — closing the fix loop tighter.

Step 5

Tool 3 needs a knowledge base

The Kodiac Agent Brand Agent is only as good as the structured content underneath it. Content is the prerequisite for a reliable Agent.

Step 6

AI infrastructure expectation

Enterprise AI teams expect RAG endpoints, vector access, semantic search testing. v2 positions Kodiac as real AI infrastructure, not just marketing tooling.

03 Application screens

Key UI views — v2 design (green identity colour)

Kodiac Content uses green (#16A34A) as its identity colour throughout. The sidebar badge, card-top accent bars, eyebrow labels, connector status dots, progress bar fills, and primary CTA button all use the green palette. The primary CTA uses .btn-primary.green (background: #16A34A).

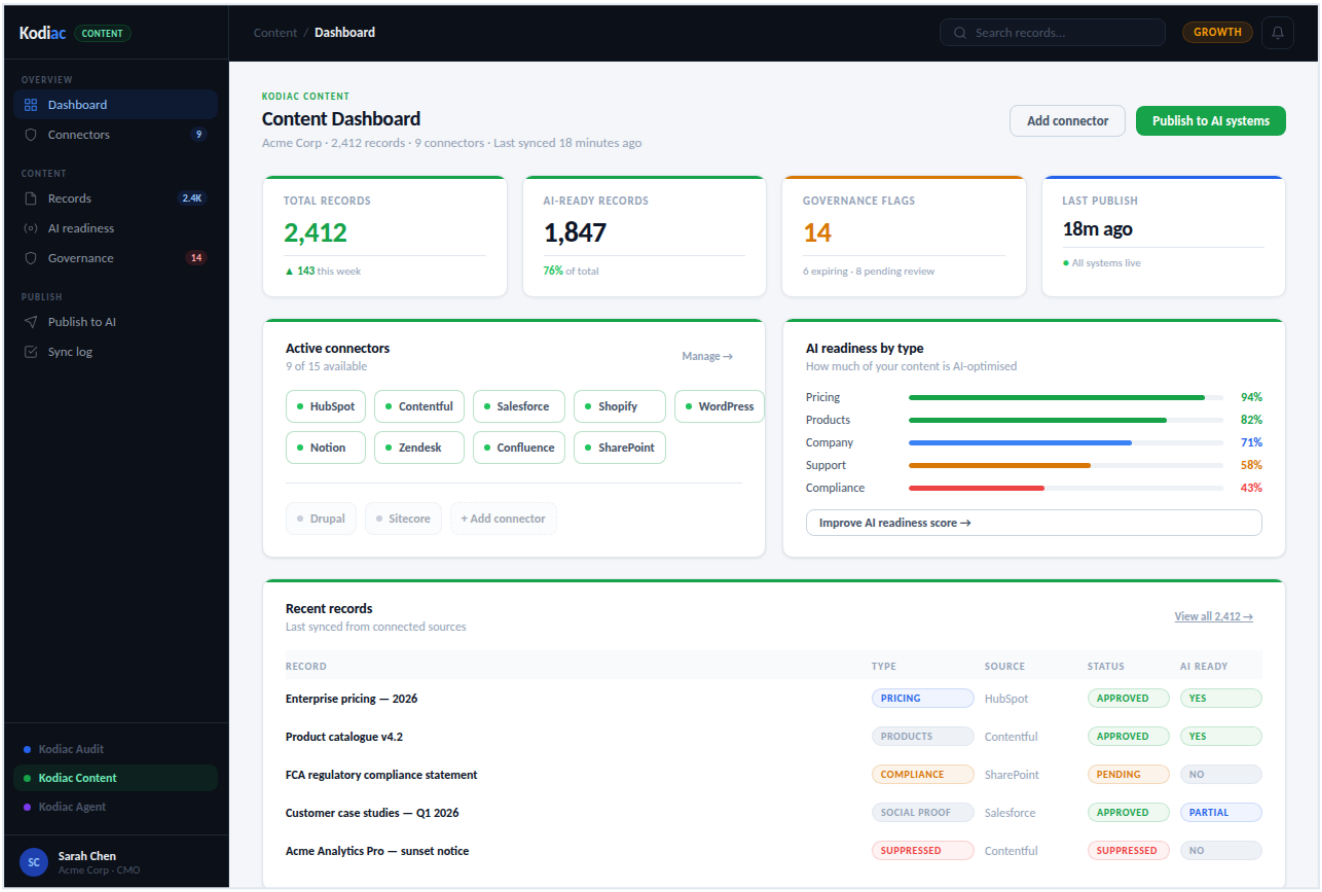

Kodiac Content — Dashboard: 9 live connector chips (green dots), AI readiness progress bars by content type, recent records table with status badges (Approved/Pending/Suppressed).

Record status badge mapping

Status	Badge class	Colour	Meaning
Approved	badge-green	Green	Content approved for AI publishing
Pending	badge-amber	Amber	Awaiting governance review before AI use
Suppressed	badge-red	Red	Actively blocked from all AI systems
Restricted	badge-dim	Grey	Not for AI use — internal only
Expiring	badge-amber	Amber	Approval expires within 30 days

04 Content Authoring — NEW

Inline structured editor with write-back to source CMS

With v1, Content was read-only: you could see records but had to fix them in the source CMS. With v2, customers can edit AI-ready records inline. The change writes back to source on save, preserving the source CMS as canonical. Optional 'Kodiac-only' override for AI-specific fields not present in source. This dramatically tightens the Audit→Content fix loop.

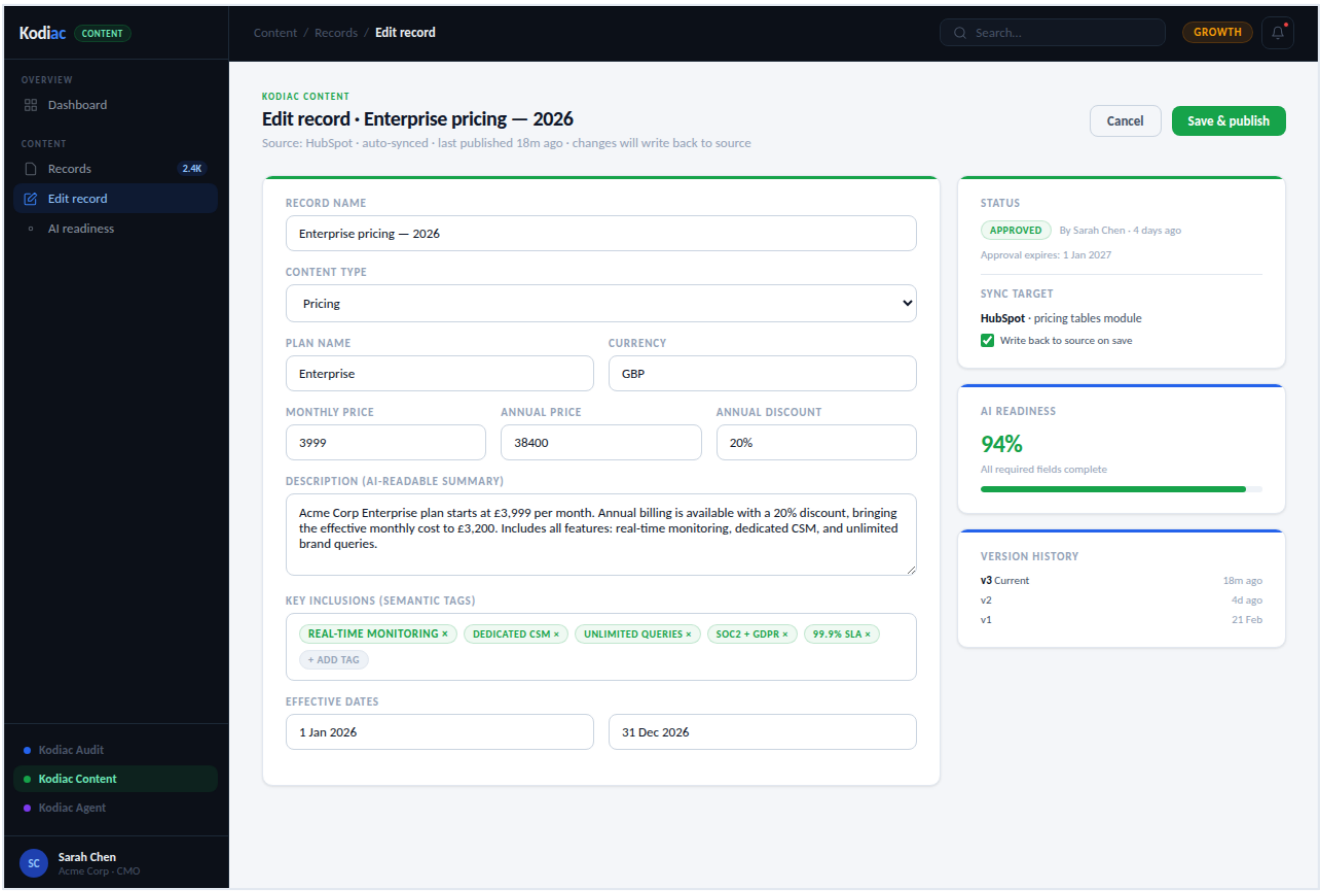

Content Authoring — structured editor for a pricing record, AI readiness score sidebar, version history, write-back-to-source toggle, semantic tag input.

Authoring feature specification

Structured field editor

Type-aware fields per content type: Pricing, Products, Company, Compliance, etc. Field validation against schema. Required-field enforcement before approval.

CORE

Semantic tag input

Free-form tag input with auto-suggest. Tags become searchable metadata and improve AI retrieval. Reused across records to build content taxonomy automatically.

CORE

Write-back to source

On save, push changes back to source CMS via connector API. Source remains canonical. Optional 'Kodiac-only' override for fields not present in source.

CORE

Live AI readiness

As user fills fields, AI readiness score updates in real time. Surfaces which missing fields would have biggest impact on AI retrieval quality.

CORE

Version history & rollback

Every save creates a version. Diff view between any two versions. Rollback restores prior version (also writes back to source).

CORE

Approval workflow integration

Editor pushes draft to approval queue (governance flag). Approved record auto-published to AI surfaces. Rejection routes back to editor with comments.

CORE

Bulk authoring mode

Edit multiple records simultaneously — useful when 20+ low-scoring pages from Audit Fix list need similar improvements (e.g. structured pricing data).

GROWTH+

05 Developer Surface — NEW

RAG endpoint, vector index, embedding config, API testing

Enterprise AI teams expect real AI infrastructure: RAG endpoints, vector store config, embedding strategy controls, semantic search testing tools. v1 mentioned 'middleware delivery API' but showed no developer surface. v2 makes Kodiac Content a first-class AI infrastructure layer that any downstream AI tool can integrate with confidently.

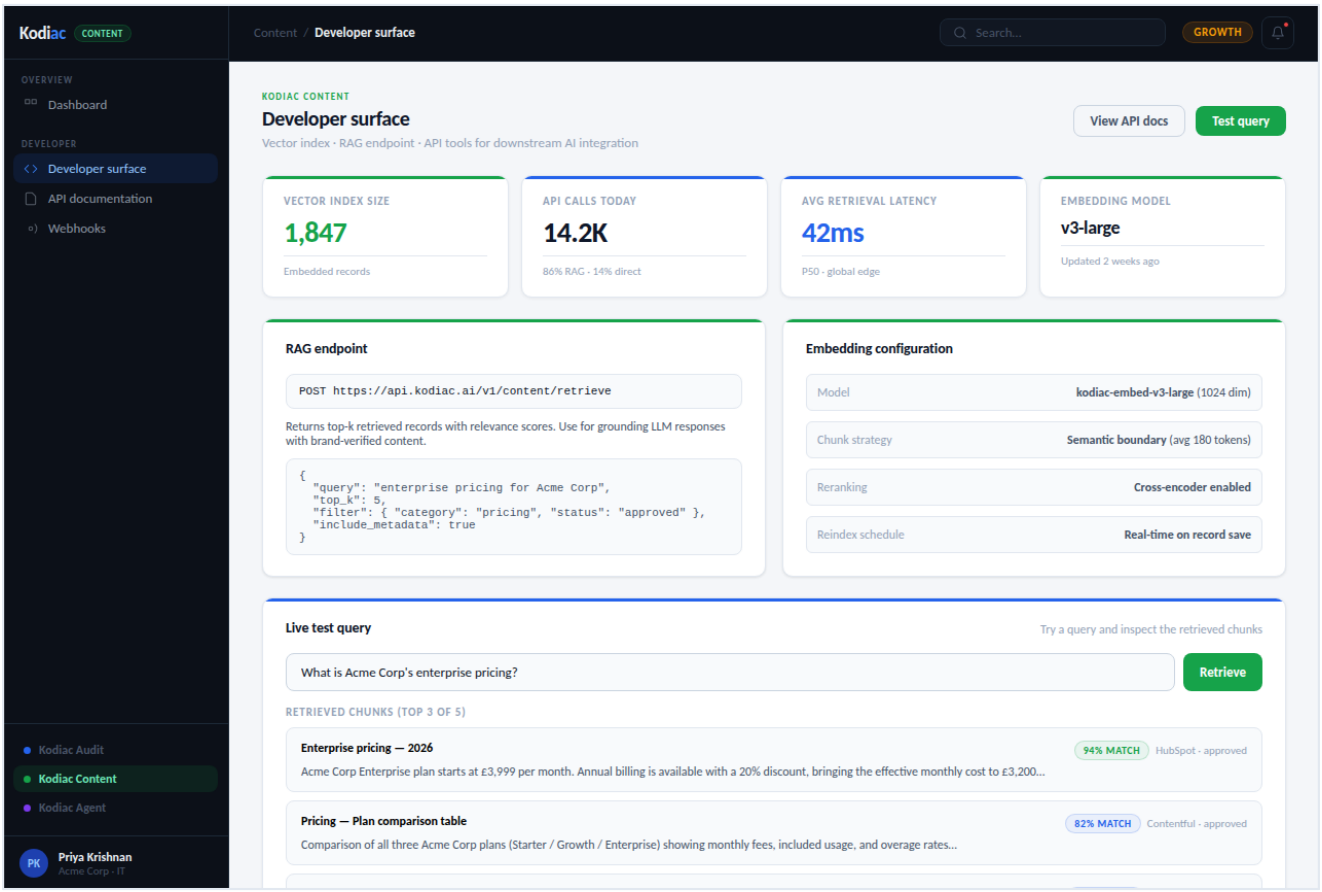

Developer surface — vector index size, RAG endpoint URL with sample request, embedding configuration (model, chunk strategy, reranking), live test query with retrieved chunks.

Developer surface feature specification

RAG retrieval endpoint

POST /v1/content/retrieve. Returns top-k records with relevance scores. Filter by category, status, locale, tag. Use to ground LLM responses with brand-verified content.

CORE

Vector index

Auto-built and maintained on every record save. Embedding model: kodiac-embed-v3-large (1024-dim). Chunk strategy: semantic boundary, avg 180 tokens.

CORE

Reranking

Cross-encoder reranking on retrieved candidates for higher accuracy on top-k results. Configurable: enable/disable, rerank-top-N, custom rerank model.

GROWTH+

Embedding config

Choose embedding model (small/large/custom). Configure chunk size, chunk strategy (semantic/fixed/hybrid), overlap percentage.

ENTERPRISE

Live test query

In-app playground to test retrieval. Enter a query, see retrieved chunks with relevance scores. Critical for AI teams validating retrieval quality before integration.

CORE

Webhooks

Subscribe to content lifecycle events: record.created, record.updated, record.published, record.suppressed. Enables downstream sync to custom AI surfaces.

GROWTH+

API rate limits

Starter: 100 rpm. Growth: 1,000 rpm. Enterprise: 10,000 rpm or custom. Per-API-key. Burst tolerance: 2x rate for 60s.

CORE

06 User journeys

Updated personas for v2 — inline authoring + dev integration

Role	Goal	Journey steps
Content Operations Lead	*Centralise, normalise and edit content for AI readiness*	**1.** Connects existing CMS sources: Sitecore, AEM, Contentful, WordPress, Drupal **2.** Adds file-based sources: SharePoint, Dropbox, Google Drive, OneDrive **3.** Maps source content structure to the unified AI-ready content model **4.** Edits low-scoring records inline in Kodiac Authoring (NEW) **5.** Bulk-updates 20 records pushed from Audit Fix list (NEW) **6.** Publishes normalised, governed content to middleware delivery layer
Enterprise Architect / IT	*Build a reliable, governable content pipeline*	**1.** Configures connector credentials and auth for each source system **2.** Sets incremental delta sync rules — avoiding expensive full re-imports **3.** Defines versioning strategy and conflict resolution across overlapping sources **4.** Configures access controls, data residency, and encryption policies **5.** Monitors pipeline health, sync failures, and content throughput **6.** Integrates content delivery API with downstream AI tools and surfaces
Brand Governance Manager	*Control what AI is permitted to access and use*	**1.** Reviews all content surfaced by normalisation engine before it goes live **2.** Applies status flags: approved, restricted, time-limited, suppressed **3.** Reviews authored drafts in approval queue (NEW — from Authoring) **4.** Sets content expiry rules — auto-archiving outdated or promotional content **5.** Reviews audit log of what AI systems retrieved and from which sources **6.** Signs off on content model updates to maintain governance standards

Role	Goal	Journey steps
AI / Product Engineering Team	*Integrate Kodiac Content as AI infrastructure (NEW)*	**1.** Queries RAG retrieval endpoint for grounded LLM responses (NEW) **2.** Uses vector index for semantic search across brand content (NEW) **3.** Configures embedding model + chunk strategy for their use case (NEW) **4.** Tests retrieval quality in live test query playground (NEW) **5.** Subscribes to webhooks for downstream content lifecycle events (NEW) **6.** Integrates Kodiac Content into custom AI agents, copilots, RAG pipelines (NEW)

07 Product feature specification

v1 capabilities (still core)

Universal content connectors

Pre-built adapters for Sitecore, AEM, Contentful, WordPress, Drupal, Sanity, SharePoint, Dropbox, Google Drive, OneDrive, Salesforce, HubSpot, Shopify, Zendesk, and more. Incremental delta sync — not expensive full re-imports. Each connector handles its own auth, pagination, rate limits, and asset resolution independently.

CORE

Normalisation engine

Transforms structurally different content from every source into a single unified, AI-readable content model. Handles field mapping, asset linking, locale resolution, version reconciliation, and deduplication across sources.

CORE

Structured content modelling

Define content models optimised for AI consumption: semantic field names, explicit content types, relationship graphs, rich taxonomy, and metadata schemas. Standard (out-of-the-box) or fully customised to brand architecture.

CORE

Governance & approval layer

Content status flags (approved, restricted, time-limited, suppressed), role-based access controls, approval workflows, and full audit trails. Precise control over what AI systems can and cannot access — with a complete log of every retrieval.

CORE

Middleware delivery API

Content delivery API that AI systems, agents, RAG pipelines, and downstream tools query to retrieve the most trusted, current version of brand content. Source attribution baked in. Filtering by content type, status, locale, topic, recency.

CORE

Structure performance monitor

Tracks which content is retrieved, by which systems, and how frequently. Surfaces content gaps that AI ignores and high-value items performing well. Feeds directly into Kodiac Audit gap analysis when connected.

ANALYTICS

Schema.org & JSON-LD export

Auto-generates and publishes structured data markup alongside content, improving how AI training pipelines, search engines and retrieval systems index your brand.

OUTPUT

Content quality scoring

Automated scoring of each content item against AI-readiness criteria: completeness, recency, semantic clarity, metadata richness, structural consistency. Surfaces content needing remediation before AI exposure.

INTELLIGENCE

Delta sync & versioning

Efficient incremental sync — only changed content is processed. Full version history maintained — enabling rollback and point-in-time reconstruction.

CORE

On-premise deployment

Containerised deployment for enterprises with data residency or sovereignty requirements that prevent cloud-hosted content processing.

ENTERPRISE

08 Pricing model

Usage metric: content records managed + API calls delivered

Kodiac Content has two independent usage levers. Records managed scales with content estate size. API calls delivered scales with how many AI surfaces query the content layer (now including RAG retrieval calls from the new developer surface). These grow at different rates: a stable content estate may plateau on records while API calls spike as more AI tools integrate with the layer.

Tier	Monthly fee	Records included	API calls included	Overage (records)	Connectors	Sync	Authoring
Starter	£799/mo	10K	50K	£40/1K records	3	Weekly	Single record
Growth	£2,499/mo	50K	500K	£30/1K records	10	Daily delta	Bulk authoring
Enterprise	£5,999/mo	250K	5M	£20/1K records	Unlimited	Real-time	Bulk + custom workflows

Annual billing: –20% on all tiers. **On-premise option:** Enterprise only, custom pricing. **No feature gates:** all tiers access core features (connectors, normalisation, governance, delivery API, authoring, RAG endpoint). Tiers differ in volumes, cadence, and advanced features (bulk authoring, custom embedding config, on-prem).

09 Go-to-market

Positioning shift, expansion motion, dependency chain

Positioning shift v1 → v2

Dimension	v1 positioning	v2 positioning
Headline	AI-ready content delivery	The complete content layer for AI
Audience	Content ops + IT	Content ops + IT + AI/Product engineering
Proof point	Connector breadth	Connector breadth + inline authoring + AI infrastructure
Selling motion	'Now you know the problem; here's how to fix it'	'Now you know the problem; fix it inline; expose it to every AI surface'

What makes Content sticky

Content estate integration creates lock-in

A customer who connects their Sitecore, AEM or SharePoint content estate is extremely unlikely to churn. The depth of integration is our primary moat.

RETENTION DRIVER

It is the prerequisite for Kodiac Agent

Content provides the knowledge base that powers the Brand Agent. Customers who want Tool 3 must have Tool 2 connected and well-structured.

RETENTION DRIVER

Inline authoring closes the Audit fix loop

Every Audit Fix-list item that lands in Content authoring tightens the cross-product workflow. The customer experiences 'I see the problem in Audit, I fix it in Content, my AI score recovers' — the platform feels like one product, not three.

RETENTION DRIVER

Developer surface creates AI engineering dependency

Once a customer's RAG pipelines, copilots, or custom agents query the Kodiac Content RAG endpoint, replacing it requires re-engineering AI infrastructure. This is a much deeper switching cost than a marketing tool would create.

RETENTION DRIVER

Links to other documents

Document	Relationship
Kodiac Audit Spec v2 (Tool 1)	The source of the fix list that drives Content adoption. Inline authoring closes the loop.
Kodiac Agent Spec v2 (Tool 3)	Tool 3 uses Content as its knowledge base. Read after this document.
Kodiac Workspace Spec	Cross-product platform layer. Workspace handles access, billing, onboarding, notifications, multi-brand.
Kodiac Master Index	Platform overview, pricing comparison, revenue model. Start here for board presentations.
App UX Spec v2.0	Green identity colour implementation, connector chip spec, authoring layout, developer surface UI.

Kodiac Content v2 · Product & GTM Spec · 2026 · Confidential